


Lecture 4

Writing Classes





Writing Classes

- **We've been using predefined classes. Now we will learn to write our own classes to define objects**
- **Chapter 4 focuses on:**
 - **class definitions**
 - **instance data**
 - **encapsulation and Java modifiers**
 - **method declaration and parameter passing**
 - **constructors**

Outline



Anatomy of a Class

Encapsulation

Anatomy of a Method



Writing Classes

- The programs we've written in previous examples have used classes defined in the Java standard class library
- Now we will begin to design programs that rely on classes that we write ourselves
- The class that contains the `main` method is just the starting point of a program
- True object-oriented programming is based on defining classes that represent objects with well-defined characteristics and functionality

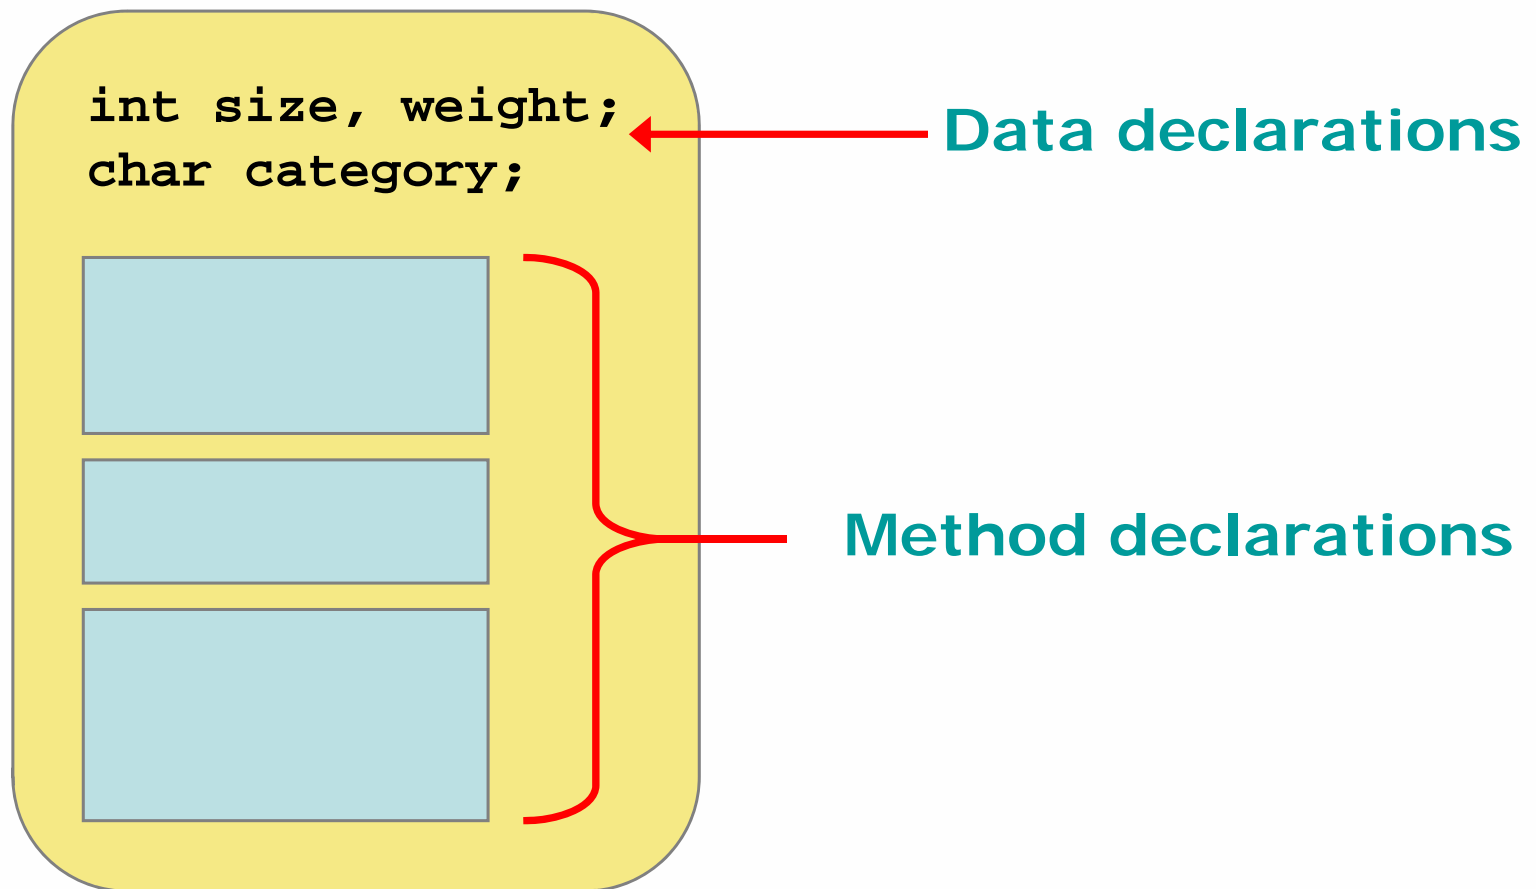


Classes and Objects

- Recall from our overview of objects in Chapter 1 that an object has *state* and *behavior*
- Consider a six-sided die (singular of dice)
 - It's state can be defined as which face is showing
 - It's primary behavior is that it can be rolled
- We can represent a die in software by designing a class called `Die` that models this state and behavior
 - The class serves as the blueprint for a die object
- We can then instantiate as many die objects as we need for any particular program

Classes

- A class can contain data declarations and method declarations





Classes

- **The values of the data define the state of an object created from the class**
- **The functionality of the methods define the behaviors of the object**
- **For our `Die` class, we might declare an integer that represents the current value showing on the face**
- **One of the methods would “roll” the die by setting that value to a random number between one and six**



Classes

- We'll want to design the `Die` class with other data and methods to make it a versatile and reusable resource
- Any given program will not necessarily use all aspects of a given class
- See [RollingDice.java](#) (page 157)
- See [Die.java](#) (page 158)



The Die Class

- The `Die` class contains two data values
 - a constant `MAX` that represents the maximum face value
 - an integer `faceValue` that represents the current face value
- The `roll` method uses the `random` method of the `Math` class to determine a new face value
- There are also methods to explicitly set and retrieve the current face value at any time



The toString Method

- All classes that represent objects should define a `toString` method
- The `toString` method returns a character string that represents the object in some way
- It is called automatically when an object is concatenated to a string or when it is passed to the `println` method



Constructors

- As mentioned previously, a *constructor* is a special method that is used to set up an object when it is initially created
- A constructor has the same name as the class
- The `Die` constructor is used to set the initial face value of each new die object to one
- We examine constructors in more detail later in this chapter



Data Scope

- The *scope* of data is the area in a program in which that data can be referenced (used)
- Data declared at the class level can be referenced by all methods in that class
- Data declared within a method can be used only in that method
- Data declared within a method is called *local data*
- In the `Die` class, the variable `result` is declared inside the `toString` method -- it is local to that method and cannot be referenced anywhere else




Instance Data

- The `faceValue` variable in the `Die` class is called *instance data* because each instance (object) that is created has its own version of it
- A class declares the type of the data, but it does not reserve any memory space for it
- Every time a `Die` object is created, a new `faceValue` variable is created as well
- The objects of a class share the method definitions, but each object has its own data space
- That's the only way two objects can have different states

Instance Data

- We can depict the two `Die` objects from the `RollingDice` program as follows:



Each object maintains its own `faceValue` variable, and thus its own state

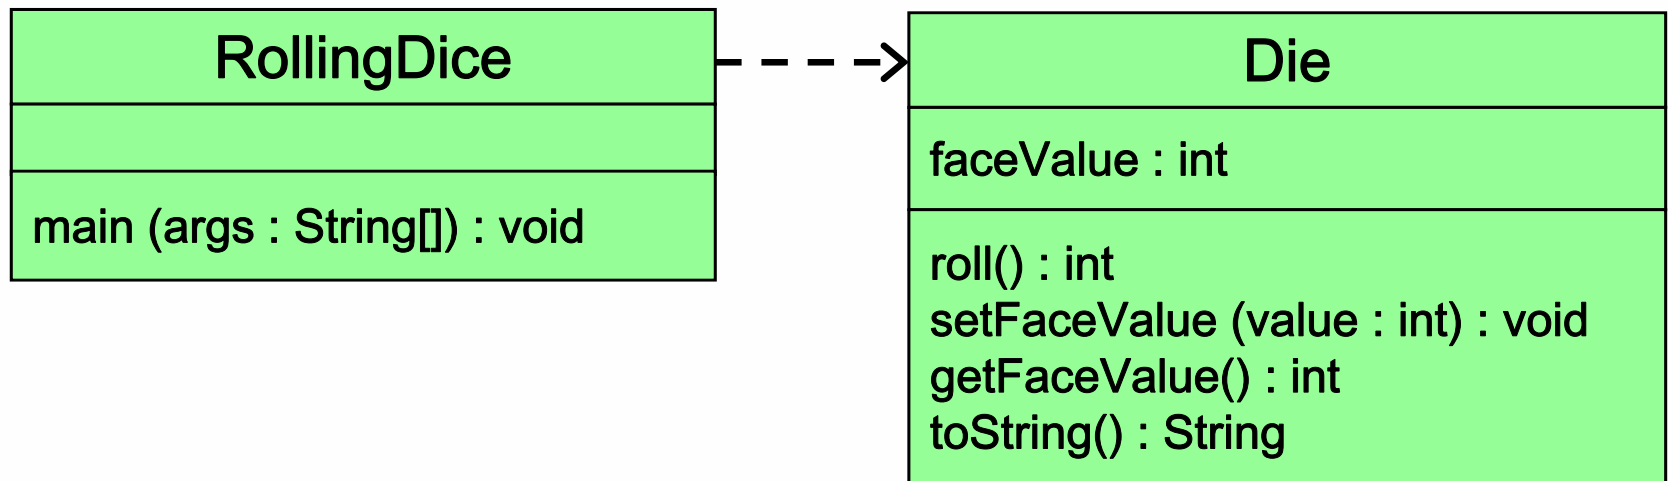


UML Diagrams

- UML stands for the *Unified Modeling Language*
- *UML diagrams* show relationships among classes and objects
- A UML *class diagram* consists of one or more classes, each with sections for the class name, attributes (data), and operations (methods)
- Lines between classes represent *associations*
- A dotted arrow shows that one class *uses* the other (calls its methods)

UML Class Diagrams

- A UML class diagram for the RollingDice program:



Outline

Anatomy of a Class



Encapsulation

Anatomy of a Method



Encapsulation

- We can take one of two views of an object:
 - internal - the details of the variables and methods of the class that defines it
 - external - the services that an object provides and how the object interacts with the rest of the system
- From the external view, an object is an *encapsulated* entity, providing a set of specific services
- These services define the *interface* to the object




Encapsulation

- One object (called the *client*) may use another object for the services it provides
- The client of an object may request its services (call its methods), but it should not have to be aware of how those services are accomplished
- Any changes to the object's state (its variables) should be made by that object's methods
- We should make it difficult, if not impossible, for a client to access an object's variables directly
- That is, an object should be *self-governing*

Encapsulation

- An encapsulated object can be thought of as a *black box* -- its inner workings are hidden from the client
- The client invokes the interface methods of the object, which manages the instance data





Visibility Modifiers

- In Java, we accomplish encapsulation through the appropriate use of *visibility modifiers*
- A *modifier* is a Java reserved word that specifies particular characteristics of a method or data
- We've used the `final` modifier to define constants
- Java has three visibility modifiers: `public`, `protected`, and `private`
- The `protected` modifier involves inheritance, which we will discuss later



Visibility Modifiers

- Members of a class that are declared with *public visibility* can be referenced anywhere
- Members of a class that are declared with *private visibility* can be referenced only within that class
- Members declared without a visibility modifier have *default visibility* and can be referenced by any class in the same package
- An overview of all Java modifiers is presented in Appendix E



Visibility Modifiers

- **Public variables violate encapsulation because they allow the client to “reach in” and modify the values directly**
- **Therefore instance variables should not be declared with public visibility**
- **It is acceptable to give a constant public visibility, which allows it to be used outside of the class**
- **Public constants do not violate encapsulation because, although the client can access it, its value cannot be changed**



Visibility Modifiers

- **Methods that provide the object's services are declared with public visibility so that they can be invoked by clients**
- **Public methods are also called *service methods***
- **A method created simply to assist a service method is called a *support method***
- **Since a support method is not intended to be called by a client, it should not be declared with public visibility**

Visibility Modifiers

	<code>public</code>	<code>private</code>
Variables	Violate encapsulation	Enforce encapsulation
Methods	Provide services to clients	Support other methods in the class



Accessors and Mutators

- Because instance data is private, a class usually provides services to access and modify data values
- An *accessor method* returns the current value of a variable
- A *mutator method* changes the value of a variable
- The names of accessor and mutator methods take the form `getX` and `setX`, respectively, where `x` is the name of the value
- They are sometimes called “getters” and “setters”



Mutator Restrictions

- The use of mutators gives the class designer the ability to restrict a client's options to modify an object's state
- A mutator is often designed so that the values of variables can be set only within particular limits
- For example, the `setFaceValue` mutator of the `Die` class should have restricted the value to the valid range (1 to `MAX`)
- We'll see in Chapter 5 how such restrictions can be implemented

Outline

Anatomy of a Class

Encapsulation



Anatomy of a Method

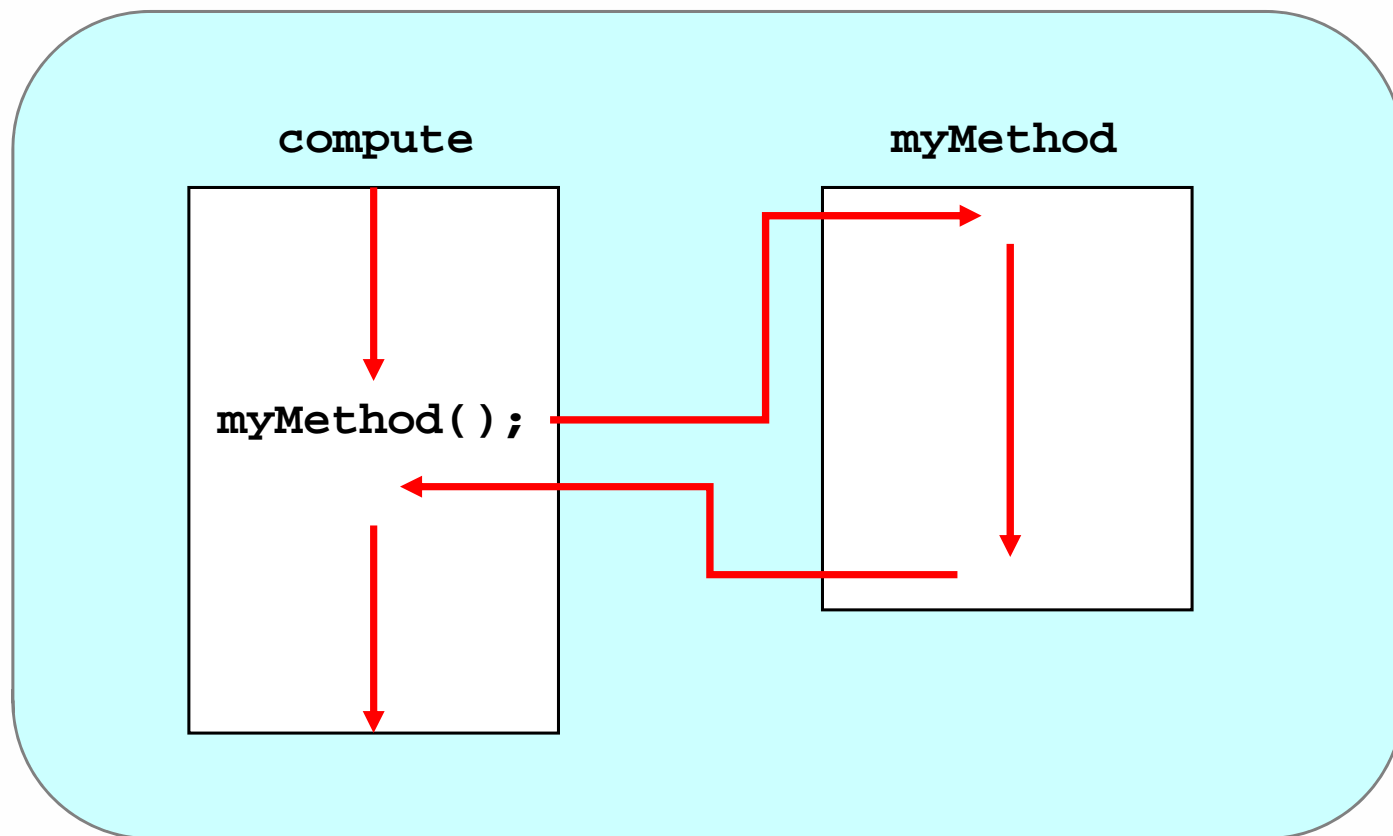


Method Declarations

- Let's now examine method declarations in more detail
- A *method declaration* specifies the code that will be executed when the method is invoked (called)
- When a method is invoked, the flow of control jumps to the method and executes its code
- When complete, the flow returns to the place where the method was called and continues
- The invocation may or may not return a value, depending on how the method is defined

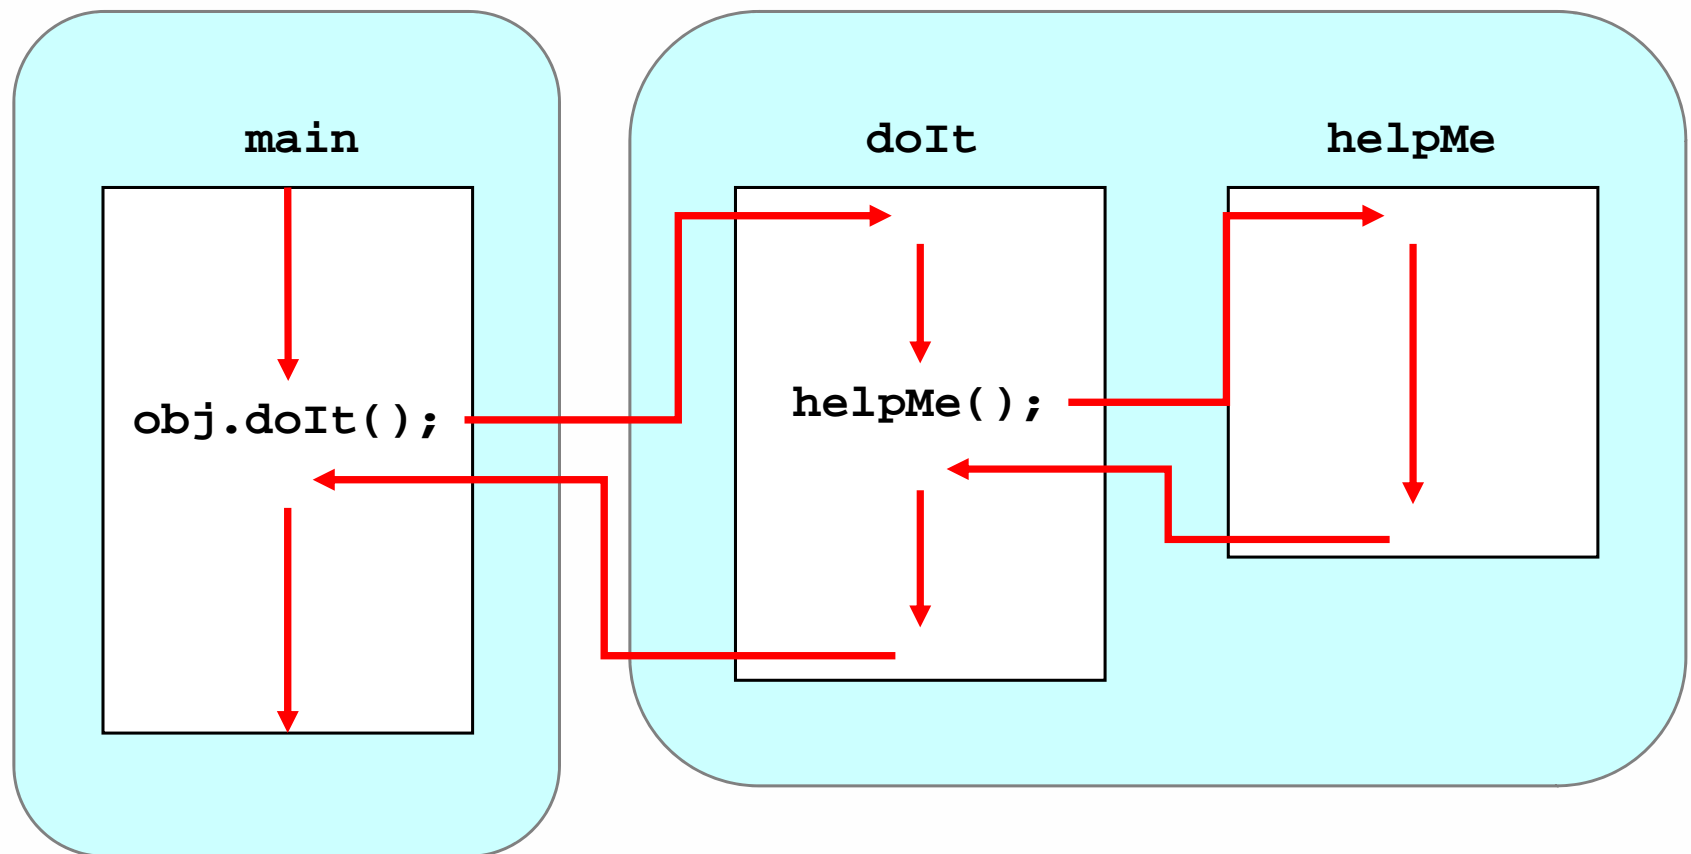
Method Control Flow

- If the called method is in the same class, only the method name is needed



Method Control Flow

- The called method is often part of another class or object



Method Header

- A method declaration begins with a *method header*

```
char calc (int num1, int num2, String message)
```

Diagram illustrating the components of a method header:

- return type**: Points to the `char` type.
- method name**: Points to the `calc` identifier.
- parameter list**: Points to the list of parameters in parentheses: `(int num1, int num2, String message)`.

The parameter list specifies the type and name of each parameter

The name of a parameter in the method declaration is called a *formal parameter*

Method Body

- The method header is followed by the *method body*

```
char calc (int num1, int num2, String message)
{
    int sum = num1 + num2;
    char result = message.charAt (sum);

    return result;
}
```

↑
The return expression
must be consistent with
the return type

sum and result
are local data

They are created
each time the
method is called, and
are destroyed when
it finishes executing



The return Statement

- The *return type* of a method indicates the type of value that the method sends back to the calling location
- A method that does not return a value has a `void` return type
- A *return statement* specifies the value that will be returned


`return expression;`

- Its expression must conform to the return type

Parameters

- When a method is called, the *actual parameters* in the invocation are copied into the *formal parameters* in the method header

```
ch = obj.calc (25, count, "Hello");
```



A diagram illustrating the mapping of actual parameters to formal parameters. A horizontal teal line separates the invocation from the method signature. Three red arrows point from the arguments in the invocation to the parameters in the signature: the first arrow points from '25' to 'num1', the second from 'count' to 'num2', and the third from '"Hello"' to 'message'.

```
char calc (int num1, int num2, String message)
{
    int sum = num1 + num2;
    char result = message.charAt (sum);

    return result;
}
```



Local Data

- As we've seen, local variables can be declared inside a method
- The formal parameters of a method create *automatic local variables* when the method is invoked
- When the method finishes, all local variables are destroyed (including the formal parameters)
- Keep in mind that instance variables, declared at the class level, exists as long as the object exists



Bank Account Example

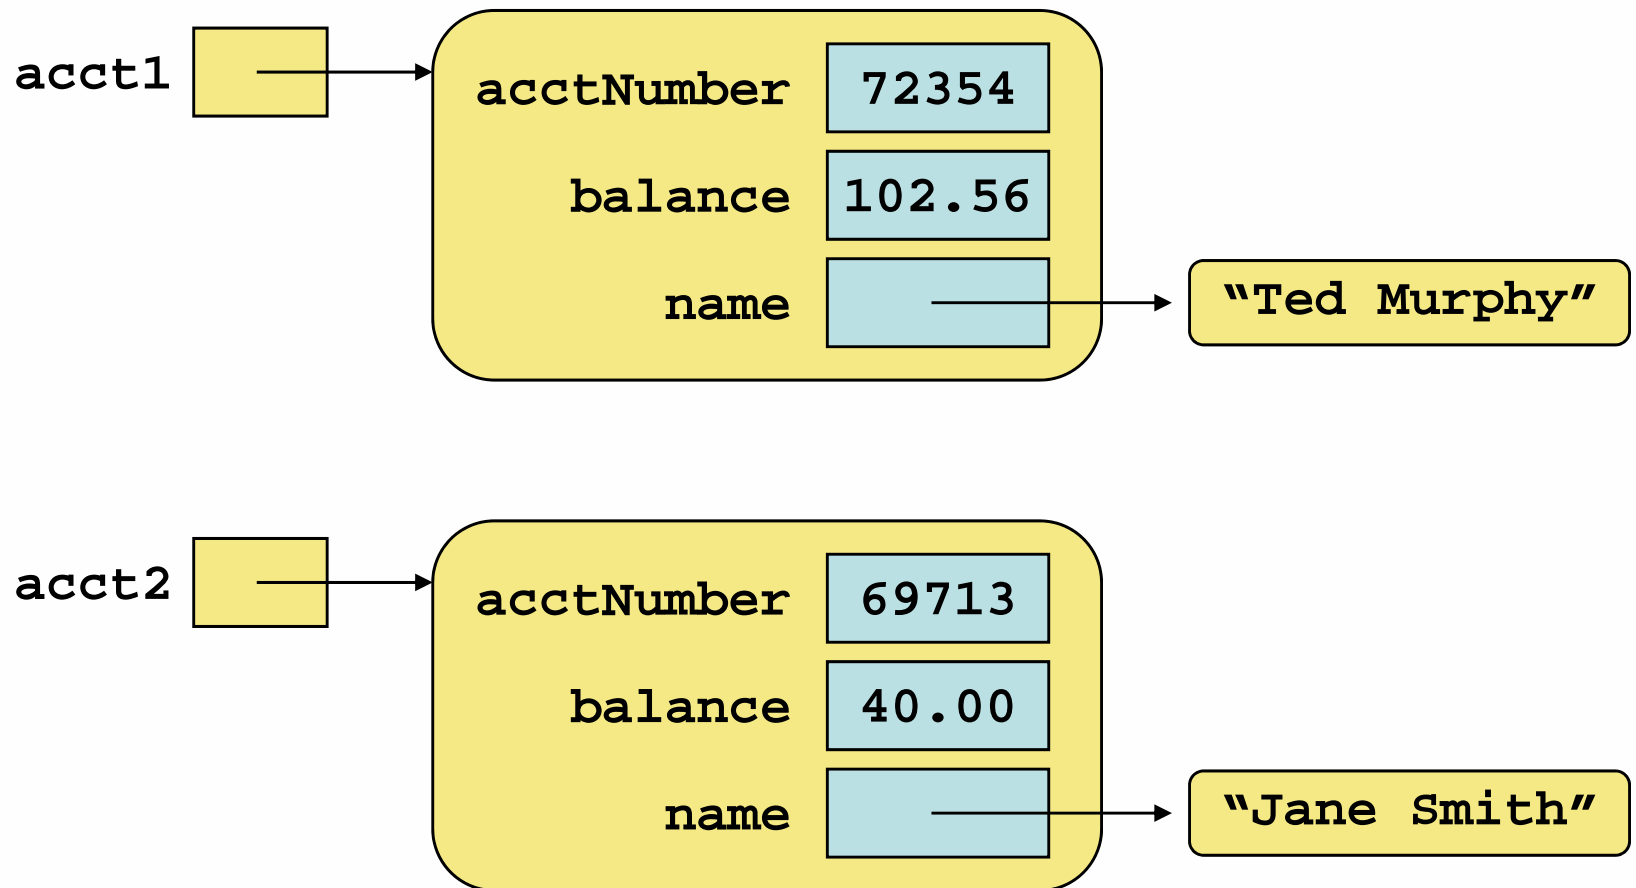
- **Let's look at another example that demonstrates the implementation details of classes and methods**
- **We'll represent a bank account by a class named Account**
- **It's state can include the account number, the current balance, and the name of the owner**
- **An account's behaviors (or services) include deposits and withdrawals, and adding interest**



Driver Programs

- A *driver program* drives the use of other, more interesting parts of a program
- Driver programs are often used to test other parts of the software
- The `Transactions` class contains a `main` method that drives the use of the `Account` class, exercising its services
- See [Transactions.java](#) (page 172)
- See [Account.java](#) (page 173)

Bank Account Example





Bank Account Example

- **There are some improvements that can be made to the Account class**
- **Formal getters and setters could have been defined for all data**
- **The design of some methods could also be more robust, such as verifying that the `amount` parameter to the `withdraw` method is positive**



Constructors Revisited

- Note that a constructor has no return type specified in the method header, not even `void`
- A common error is to put a return type on a constructor, which makes it a “regular” method that happens to have the same name as the class
- The programmer does not have to define a constructor for a class
- Each class has a *default constructor* that accepts no parameters



Summary

- **Lecture 4 focused on:**
 - **class definitions**
 - **instance data**
 - **encapsulation and Java modifiers**
 - **method declaration and parameter passing**
 - **constructors**